

AJAY KUMAR

+91-6397053766 | a.kumar01c@gmail.com | [LinkedIn](#) | [GitHub](#)

Skills

Data Engineering & ML – Databricks, Delta Lake, DBT, PySpark, Snowflake, Airflow, Kafka, Medallion Architecture

Big Data & Databases – Apache Spark, Hadoop, MySQL, PostgreSQL, Cassandra, MongoDB

DevOps & Cloud – Docker, Linux, Azure (ADF, ADLS Gen2), AWS (Athena, Redshift, S3, Glue)

Programming & Analytics – Python (pandas, NumPy for statistical analysis), SQL, C, Power BI, IBM Cognos Analytics

Experience

DAETL Solutions – Data Engineer

Jun'25 – Mar'26 | Remote

- Migrated legacy Informatica ETL workflows (**~100 GB/day**) to Azure Databricks using PySpark, Spark SQL, and Delta Lake on ADLS Gen2 with Medallion architecture, improving pipeline performance **by 40%**.
- Owned development of a **data quality validation framework in Databricks covering 50+ Silver-layer tables**, automating row-count reconciliation, column-level checks, and SQL-based business rule validation to catch mismatches before downstream reporting.
- Partnered with QA analysts and business users to translate Excel mapping rules into ADF-orchestrated Databricks pipelines, tuning Spark jobs to **reduce runtime by 50%**.

VITISPRO LLC – Data Engineer

Apr'23 – Nov'23 | Remote

- Owned migration of key data functions and databases from GCP to AWS, aligning with engineering stakeholders on data cleaning rules and schema standards across environments; **cut project expenses by 55%**.
- Led cross-functional automation of Eurostat and Comtrade API ingestion, coordinating with analytics users to define market and economic dataset requirements; **reduced data retrieval time by 40%**.
- Led enhancement of **data ingestion workflows to parse 5 formats**, including PNG and SVG, enabling internal teams to analyze heterogeneous data through a standardized pipeline.

Projects

CityFlux: Real-Time Data Streaming Pipeline (Kafka & Spark) [\[GITHUB\]](#)

- Designed and deployed a scalable real-time data pipeline using Apache Kafka and Spark to ingest and process streaming vehicle, GPS, and weather data, achieving low-latency event processing and **reducing decision-making lag by 40%**.
- Integrated AWS services (Glue, Athena, Redshift) for automated ETL, Metadata Cataloging, and fast querying; **improved analytical reporting performance by ~35%**.
- Containerized Spark and Kafka environments using Docker and Docker Compose, **cutting down environment setup time by over 50%** and accelerating development workflows for new contributors.

Pipeplatter: Food Delivery Data Platform (Databricks, DBT, Airflow) [\[GITHUB\]](#) [\[DEMO\]](#)

- Built an end-to-end food delivery analytics platform on Databricks, modeling **7 business domains with a 25-model DBT medallion warehouse** for TB-scale batch processing.
- Automated a daily Airflow pipeline with **12+ tasks for parallel ingestion**, staged DBT runs, data quality tests, and AI enrichment; improved reliability with idempotent loading and incremental fact models.
- Deployed an intelligence layer using Streamlit, enabling business KPI communication across dashboards, natural language SQL querying, and RAG review search, which **enhance data accessibility for 5 gold marts and top-k semantic retrieval**.

Education

Hindustan College of Science and Technology

B. Tech – Computer Science Engineering

Mathura, UP

Jun'22

Courses & Certifications

- AWS Cloud Practitioner Essentials – Amazon Web Services
- IBM Data Engineering Professional Certificate, IBM Data Warehousing Certificate – IBM, Coursera
- CS50's Introduction to Databases with SQL – CS50, Cambridge, Massachusetts
- Hands On Machine Learning Course – Udemy
- Programming for Everybody – University of Michigan, Coursera
- LangChain and LangGraph for Generative AI – CampusX, FreeCodecamp