

AJAY KUMAR

+91-6397053766 | a.kumar01c@gmail.com | [LinkedIn](#) | [GitHub](#)

Skills

Data Engineering & ML – Databricks, Delta Lake, PySpark, Snowflake, Airflow, Kafka, Medallion Architecture

Big Data & Databases – Apache Spark, Hadoop, MySQL, PostgreSQL, Cassandra, MongoDB

DevOps & Cloud – Docker, Azure (ADF, ADLS Gen2), AWS (Athena, Redshift, S3, Glue)

Programming & Analytics – Python, SQL, C, SAS, Power BI, IBM Cognos Analytics

Experience

DAETL Solutions – Data Engineer

Jun'25 – Mar'26 | Remote

- Migrated legacy Informatica workflows into scalable Azure Databricks pipelines using PySpark, Spark SQL, and Delta Lake on ADLS Gen2, **implementing Medallion architecture**.
- Orchestrated **end-to-end data pipelines** using Azure Data Factory (ADF) and optimized Spark performance through partitioning, caching, and join strategies.
- Developed **data validation framework** in Databricks for row-count reconciliation, column-level checks, and business rule validation to ensure data integrity.

VITISPRO LLC – Data Engineer

Apr'23 – Nov'23 | Remote

- Optimized operational efficiency** by migrating key functions from GCP to AWS, **achieving a 55% reduction in project expenses** and streamlining cloud resource management.
- Led the **automation of bulk data collection** for Eurostat and Comtrade via API integrations, improving process efficiency and **achieving a 40% faster retrieval speed**, boosting team productivity with quicker access to critical datasets.
- Collaborated across teams for the **expansion of image extraction functionality** to support 5 file types, including PNG and SVG, enhancing data availability for analysis within internal systems.

Projects

CityFlux: Real-Time Event Stream (using Kafka & Spark) [\[LINK\]](#)

- Designed and deployed a scalable real-time data pipeline using Apache Kafka and Spark to ingest and process streaming vehicle, GPS, and weather data, achieving low-latency event processing and **reducing decision-making lag by 40%**.
- Integrated AWS services (Glue, Athena, Redshift) for automated ETL, Metadata Cataloging, and fast querying; **improved analytical reporting performance by ~35%**.
- Containerized Spark and Kafka environments using Docker and Docker Compose, **cutting down environment setup time by over 50%** and accelerating development workflows for new contributors.

RedditFlow: Scalable ETL Pipeline (using Airflow & Spark) [\[LINK\]](#)

- Built and orchestrated a production-grade ETL pipeline using Apache Airflow and PySpark to extract Reddit post data, **achieving a 30% increase in data ingestion throughput**.
- Automated end-to-end data flow across AWS S3, Glue, and Redshift, **reducing manual data handling time by over 15 hours per week** and improving pipeline reliability.
- Deployed a modular, cloud-native architecture leveraging Athena and Redshift to enable efficient querying and scalable analytics on large Reddit datasets.

Education

Hindustan College of Science and Technology

Mathura, UP

B. Tech – Computer Science Engineering

Jun'22

Courses & Certifications

- AWS Cloud Practitioner Essentials – Amazon Web Services
- IBM Data Engineering Professional Certificate, IBM Data Warehousing Certificate – IBM, Coursera
- CS50's Introduction to Databases with SQL – CS50, Cambridge, Massachusetts
- Hands On Machine Learning Course – Udemy
- Programming for Everybody – University of Michigan, Coursera
- LangChain and LangGraph for Generative AI – CampusX, FreeCodecamp